✓ 100 XP ▶

How Kubernetes deployments work

10 minutes

The drone tracking application has the following components that are deployed separately from each other. It's your job to configure deployments for these components on the cluster.

- Public tracking website
- Private in-memory cache service
- Public RESTful API
- Private persisted queue
- Private data processing service
- Private NoSQL database

Pod deployment options

There are several options to manage the deployment of pods in a Kubernetes cluster when you're using `kubectl`. You can use any of four object type definitions to deploy a pod or pods. These files make use of YAML to describe the intended state of the pod or pods that will be deployed.

Pod templates

You use a pod template to deploy pods manually. Keep in mind that a manually deployed pod isn't relaunched after it fails, is deleted, or is terminated.

Replication controllers

A replication controller uses pod templates and defines a specified number of pods that must run. The controller helps you run multiple instances of the same pod and ensures that the specified number of pods are always running on one or more nodes in the cluster. A replication controller will replace pods launched in this way with new pods if they fail, are deleted, or are terminated.

For example, assume that you deploy the drone tracking front-end website and users start accessing the website. If all the pods fail for any reason, the website is unavailable to your users unless you launch new pods. A replication controller helps you make sure that your website is always available.

Replica sets

Replica sets replace the replication controller as the preferred way to deploy replicas. A replica set includes the same functionality as a replication controller. However, it has an extra configuration option to include a selector value.

A selector allows the replica set to identify all the pods running underneath it. This feature allows you to manage pods labeled with the same value as the selector value, but not created with the replicated set.

Deployments

A deployment creates a management object one level higher than a replica set. You can use deployment management updates to manage how you update pods in a cluster.

Assume that you have five instances of your application deployed in your cluster. There are five pods running version 1.0.0 of your application.

Kubernetes Node

v1 Pod	v1 Pod	v1 Pod	v1Pod	v1 Pod
app = front-end-nginx version=1.0.0	app = front-end-nginx version=1.0.0	app = front-end-nginx version=1.0.0	app = front-end-nginx version=1.0.0	app = front-end-nginx version=1.0.0
Website	Website	Website	Website	Website

If you decide to update your application manually, you can terminate all pods and then launch new pods running version 2.0.0 of your application. With this strategy, your application will experience downtime.

What you instead want to do is execute a rolling update, where you launch pods running the new version of your application before you terminate the pods running the older version of your application. Rolling updates will launch one pod at a time instead of taking down all the older pods at once. Deployments honor the number of replicas configured in the section that describes information about replica sets. It will maintain the number of pods specified in the replica set as it terminates old pods and launches new pods.

Kubernetes Node

v1 Pod app = front-end-nginx version=1.0.0 Website	**v1 Pod** app = front-end-nginx version=1.0.0 Website	**v2 Pod** app = front-end-nginx version=2.0.0 Website	**v2Pod** app = front-end-nginx version=2.0.0 Website	**v2 Pod** app = front-end-nginx version=2.0.0 Website

Deployments, by default, provide a rolling update strategy for updating pods. You also have the option to use a re-create strategy. This strategy will terminate pods before launching new pods.

Deployments also provide you with a rollback strategy, which you can execute by using `kubectl`.

Deployments make use of YAML-based definition files and make it easy to manage deployments. Keep in mind that deployments allow you to apply any changes to your cluster. For example, you can deploy new versions of an app, update labels, and run other replicas of your pods.

`kubectl` has convenient syntax to create a deployment automatically when you're using the `kubectl run` command to deploy a pod. This command creates a deployment with the required replica set and pods. However, the command doesn't create a definition file. A best practice is to manage all deployments with deployment definition files and track changes by using a version control system.

Deployment considerations

Kubernetes has specific requirements on how you configure networking and storage for a cluster. How you configure these two aspects affects your decisions on how to expose your applications on the cluster network and store data.
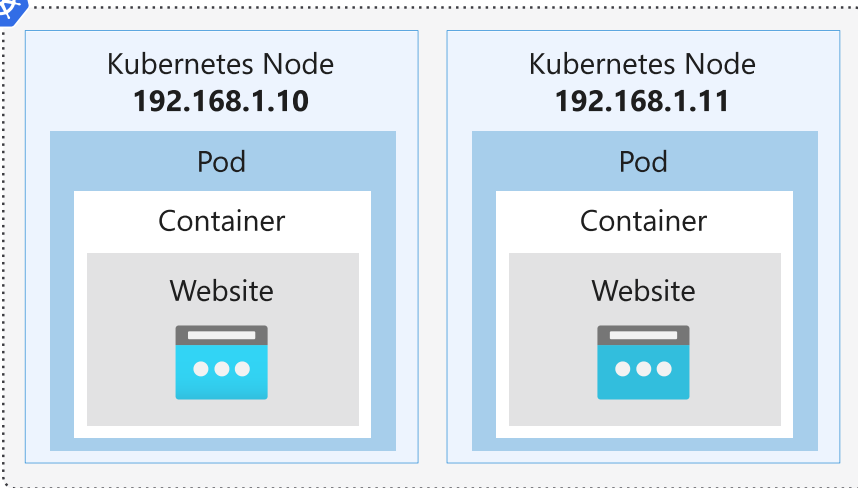
For example, each of the services in the drone tracking application has specific requirements for user access, inter-process network access, and data storage. Let's have a look at these aspects of a Kubernetes cluster and how they affect the deployment of applications.

Kubernetes networking

Assume that you have a cluster with one master and two nodes. When you add nodes to Kubernetes, an IP address is automatically assigned to each node from an internal private network range. For example, assume that your local network range is 192.168.1.0/24.
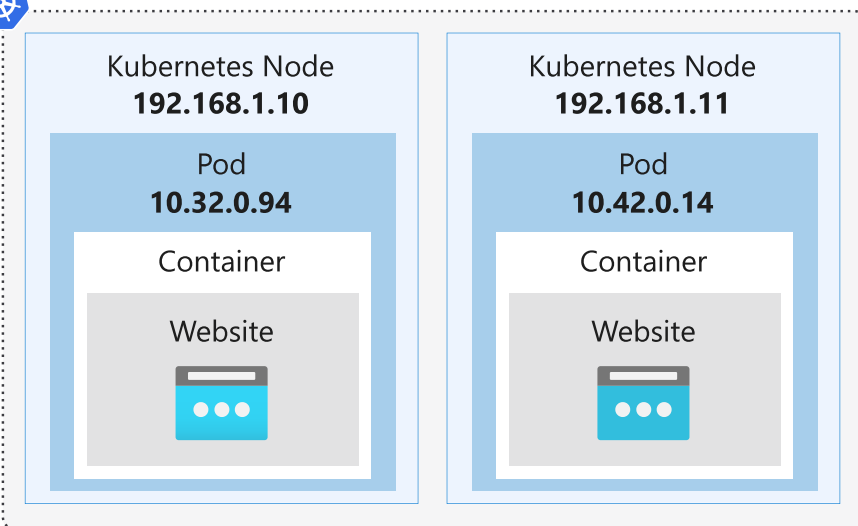
Kubernetes Cluster

Kubernetes Node **192.168.1.10**	Kubernetes Node **192.168.1.11**
Pod	Pod
Container	Container
Website	Website

Each pod that you deploy gets assigned an IP from a pool of IP addresses. For example, assume that your configuration uses the 10.32.0.0/12 network range.

Kubernetes Cluster

Kubernetes Node **192.168.1.10**	Kubernetes Node **192.168.1.11**
Pod **10.32.0.94**	Pod **10.42.0.14**
Container	Container
Website	Website

By default, the pods and nodes can't communicate with each other by using different IP address ranges.

To further complicate matters, recall that pods are transient. The pod's IP address is temporary and can't be used to reconnect to a newly created pod. This configuration affects how your application communicates with its internal components and how you and services interact with it externally.

To simplify communication, Kubernetes expects you to configure networking in such a way that:

- Pods can communicate with one another across nodes without Network Address Translation (NAT).
- Nodes can communicate with all pods, and the other way around, without NAT.

- Agents on a node can communicate with all nodes and pods.

Kubernetes offers several networking options that you can install to configure networking. Examples include Antrea, Cisco Application Centric Infrastructure (ACI), Cilium, Flannel, Kubenet, VMware NSX-T, and Weave Net.

Cloud providers also provide their networking solutions. For example, Azure Kubernetes Service (AKS) supports the Azure Virtual Network container network interface (CNI), Kubenet, Flannel, Cilium, and Antrea.

Services

A service is a Kubernetes object that provides stable networking for pods. A Kubernetes service enables communication between nodes, pods, and users of your application, both internal and external, to the cluster.

Kubernetes assigns a service an IP address on creation, just like a node or pod. These addresses get assigned from a service cluster's IP range. An example is 10.96.0.0/12. A service is also assigned a DNS name based on the service name, and an IP port.

In the drone tracking application, network communication is as follows:

- The website and RESTful API are accessible to users outside the cluster.

- The in-memory cache and message queue services are accessible to the front end and the RESTful API, respectively, but not to external users.

- The message queue needs access to the data processing service, but not to external users.

- The NoSQL database is accessible to the in-memory cache and data processing service, but not to external users.

To support these scenarios, you can configure three types of services to expose your app's components.

ClusterIP	The address assigned to a service that makes the service available to a set of services inside the cluster. For example, communication between the front-end and back-end components of your application.
NodePort	The node port, between 30000 and 32767, that the Kubernetes control plane assigns to the service. An example is 192.169.1.11 on clusters01. You then configure the service with a target port on the pod that you want to expose. For example, configure port 80 on the pod running one of the front ends. You can now access the front end through a node IP and port address.

| **LoadBalancer** | The load balancer that allows for the distribution of load between nodes running your application and exposing the pod to public network access. You typically configure load balancers when you use cloud providers. In this case, traffic from the external load balancer is directed to the pods running your application. |

In the drone tracking application, you might decide to expose the tracking website and the RESTful API by using LoadBalancer services and the data processing service with a ClusterIP.

How to group pods

Pod IP addresses change as controllers re-create them, and you might have any number of pods running. Managing pods by IP address isn't practical.

A service object allows you to target and manage specific pods in your cluster by using selector labels. You set the selector label in a service definition to match the pod label defined in the pod's definition file.

Cluster Node - 192.168.1.10
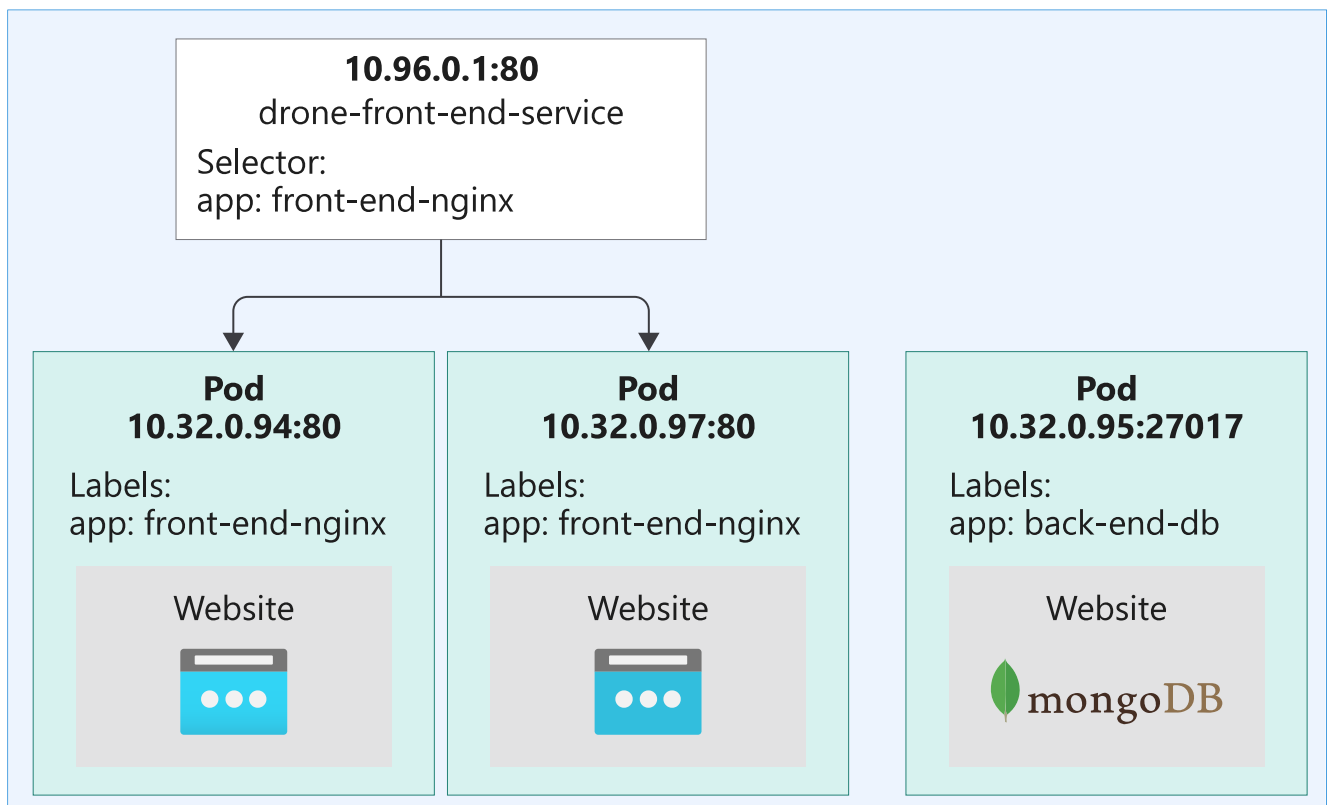
For example, assume that you have many running pods. Only a few of these pods are on the front end, and you want to set a LoadBalancer service that targets only the front-end pods. You can apply your service to expose these pods by referencing the pod label as a selector value in the service's definition file. The service will now group only the pods that match the

label. If a pod is removed and re-created, the new pod is automatically added to the service group through its matching label.

Kubernetes storage

Kubernetes uses the same storage volume concept that you find when using Docker. Docker volumes are less managed than the Kubernetes volumes because Docker volume lifetimes aren't managed. The Kubernetes volume's lifetime is an explicit lifetime that matches the pod's lifetime. This lifetime match means a volume outlives the containers that run in the pod. However, if the pod is removed, so is the volume.

Kubernetes provides options to provision persistent storage with the use of *PersistentVolumes*. You can also request specific storage for pods by using *PersistentVolumeClaims*.

Keep both of these options in mind when you're deploying application components that require persisted storage like message queues and databases.

Cloud integration considerations

In a cloud environment such as Azure, you can use several services outside the Kubernetes cluster. Recall from earlier that Kubernetes doesn't provide middleware, data-processing frameworks, databases, caches, or cluster storage systems.

In the drone tracking solution, three services provide middleware functionality that needs consideration. There's a NoSQL database, an in-memory cache service, and a message queue. For example, you might choose to use MongoDB Atlas for the NoSQL solution, Redis to manage in-memory cache and RabbitMQ, or Kafka, depending on your message queue needs.

When you're using a cloud environment such as Azure, it's a best practice to make use of services outside the Kubernetes cluster. This decision can simplify the cluster's configuration and management. For example, you can use *Azure Cache for Redis* for the in-memory caching services, *Azure Service Bus messaging* for the message queue, and *Azure Cosmos DB* for the NoSQL database.

Next unit: Exercise - Explore the functionality of a Kubernetes cluster

Continue >